Vikram Srinivasan

571-882-3057 | viksrin@outlook.com

vik-srinivasan.github.io/personal-website | linkedin.com/in/vikram-srin | github.com/vik-srinivasan

EDUCATION

Stanford University — Palo Alto, CA
Master of Science in Computer Science, *Concentration in Artificial Intelligence* — *Jan. 2025 - June 2026*
Bachelor of Science in Symbolic Systems, *Concentration in Natural Language* — *Sep. 2021 - Dec. 2025*

EXPERIENCE

Cartesia AI — San Francisco, CA
Member of Technical Staff Intern | *Python, SQL, Databricks, JavaScript, PyTorch* — *Dec 2025 - Current*
- Developed and shipped full-stack GTM engineering features, including ML-driven lead-discovery pipelines, data ingestion services, and internal dashboards that surfaced high-value companies for streamlined GTM workflows.
- Engineered backend components for customer-identification systems, building ranking models, enrichment pipelines, and prospect-scoring tools that help the GTM team prioritize the most relevant leads at scale.

CoStar Group — Washington, D.C.
Machine Learning Intern | *Python, PyTorch, AWS, Streamlit, Databricks, NoSQL* — *June 2025 - Sep. 2025*
- Built and deployed a large-scale computer vision pipeline on AWS for detecting kitchen features in real estate imagery, training and fine-tuning CNN and Vision Transformer (ViT) models in PyTorch to achieve over 90% average precision and 80% average recall on key object classes.
- Applied feature engineering, hyperparameter tuning, and iterative evaluation to improve ranking performance of detected features for downstream search and personalization systems.
- Delivered a real-time, human-in-the-loop validation dashboard in Streamlit, enabling interactive model inspection.

Software Engineering Intern | *Python, SQL, TypeScript, PyTorch, AWS, Git* — *June 2024 - Sep. 2024*
- Developed a multimodal search engine for 100K+ condo buildings, integrating text embeddings (NLP) and visual embeddings (CV) to improve retrieval accuracy and personalization.
- Built scalable AWS-based APIs for generating and serving vector representations, leveraging OpenSearch for ranking optimization and distributed query execution.
- Designed a responsive interface for interactive content discovery, enabling seamless exploration of multimodal search results.

BSE Global (Brooklyn Nets) — Brooklyn, NY
Data Analytics & Insights Intern | *Python, MySQL, Tableau* — *June 2023 - Sep 2023*
- Built a Machine Learning model using scikit-learn to classify customer segments, resulting in targeted marketing strategies and improved customer acquisition rates.
- Developed a Natural Language Processing (NLP) algorithm to analyze sentiment in sales rep-customer phone calls, aiming to refine sales strategies and performance.

PROJECTS

Geovision: Image Geolocation with ViTs | github.com/vik-srinivasan/GeoGuessrCV | *Python, PyTorch*
- Fine-tuned ViT-B/16 and ResNet-50 on a 104-class country-level geolocation task using GeoGuessr images, achieving 88.3% Top-5 accuracy, surpassing human baselines by 3–5$\times$ and outperforming GPT-4o.

Context-Aware Voice Dictation for Accessibility | github.com/vik-srinivasan/Audio-Feedback | *Python*
- Developed a Voice Dictation Tool leveraging Named Entity Recognition (NER), real-time error correction, and a lightweight memory mechanism to enhance transcription accuracy for blind and low-vision users.
- Built backend functionality in Python to integrate speech-to-text (OpenAI Whisper) and NER (using Google Gemini) for proper noun detection and correction.

RepWise | github.com/vik-srinivasan/RepWise | *TypeScript, React Native, CSS, Supabase*
- Designed and developed a React Native mobile application that generates personalized workout plans using AI, integrating the Gemini API to tailor workouts based on user preferences, workout history, and fitness goals.

SKILLS

- **Technologies**: Python, Java, JavaScript, TypeScript, SQL, PyTorch, React, Next.js, Node.js, Git, Docker, AWS
- **Expertise**: ML/Data Science, NLP, CV, Full-Stack/Front-End Development, Data Structures & Algorithms